



Container Optimization

Logistics



- **Class Hours:**
- Instructor will provide class start and end times.
- Breaks throughout class



- **Telecommunication:**
- Turn off or set electronic devices to vibrate
- Reading or attending to devices can be distracting to other students

Course Objectives



- Explain the benefits of using containers
- Understand Container images
- Optimize container images
- Container resource tuning
- Utilizing Open-Source tools

Instructor

Rod Davison

50 years experience

Academia (math, linguistics, cognitive science)

Artificial Intelligence R and D

Software Development

Data Analytics – Social Research

Project Manager

Quality and Testing

Business Analysis

Consulting and Training




Introductions

Hello!

- Name
- Job Role
- Your experience with (scale 1 - 5)
 - Docker/Containers
- Expectations for the course (please be specific)

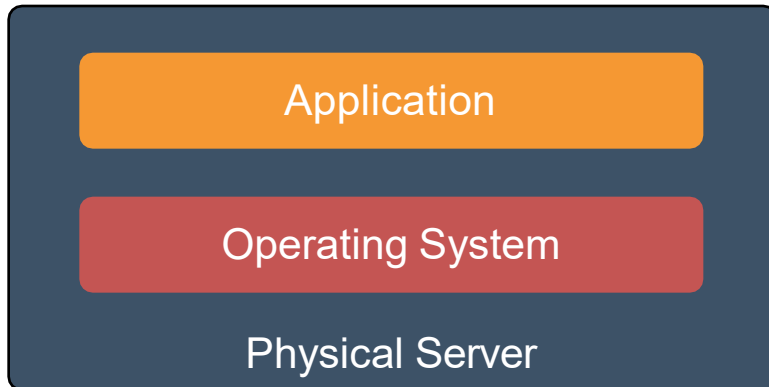
Data Center Evolution



Monolithic

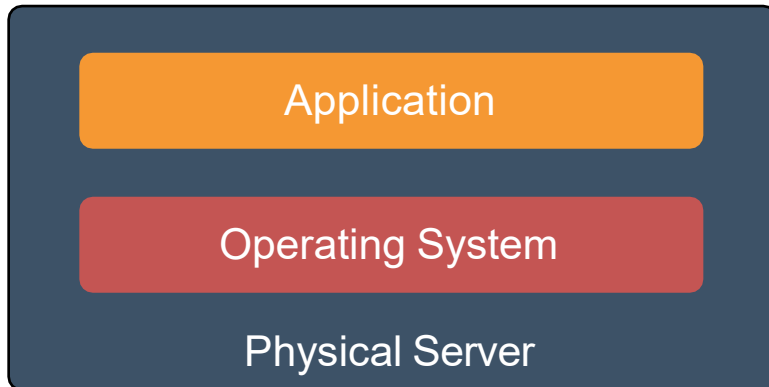


Monolithic Server Architecture



One physical server, one application

Monolithic Server Architecture



One physical server, one application

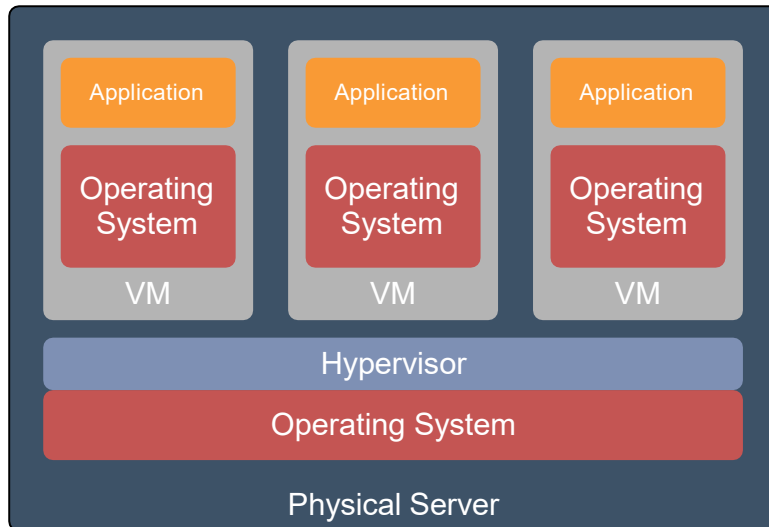
Problems

- Slow deployment times
- Cost
- Wasted resources
- Difficult to scale
- Difficult to migrate

Virtualized



Virtualized Infrastructure




One physical server, multiple applications

Discussion

What are some of the advantages and disadvantages of Virtual Machines?

Virtualized Infrastructure - Advantages

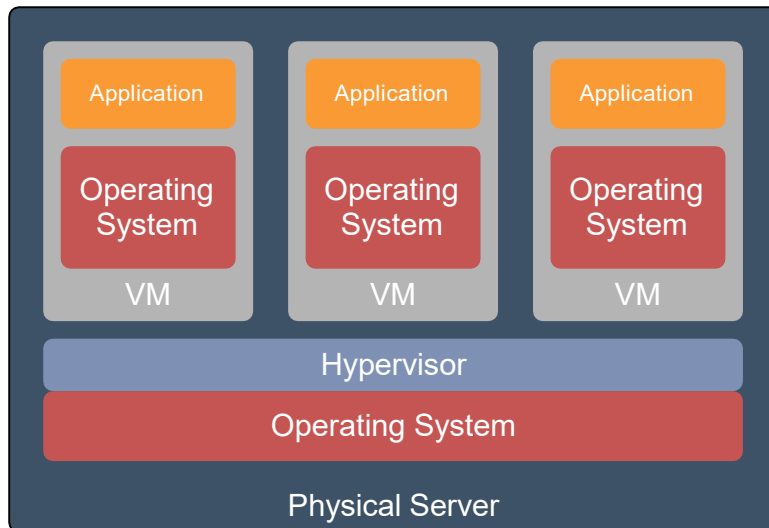



One physical server, multiple applications

Advantages

- Better resource pooling
- Easier to Scale
- Enables Cloud/IaaS
 - Rapid elasticity
 - Pay as you go model

Virtualized Infrastructure - Limitations



One physical server, multiple applications

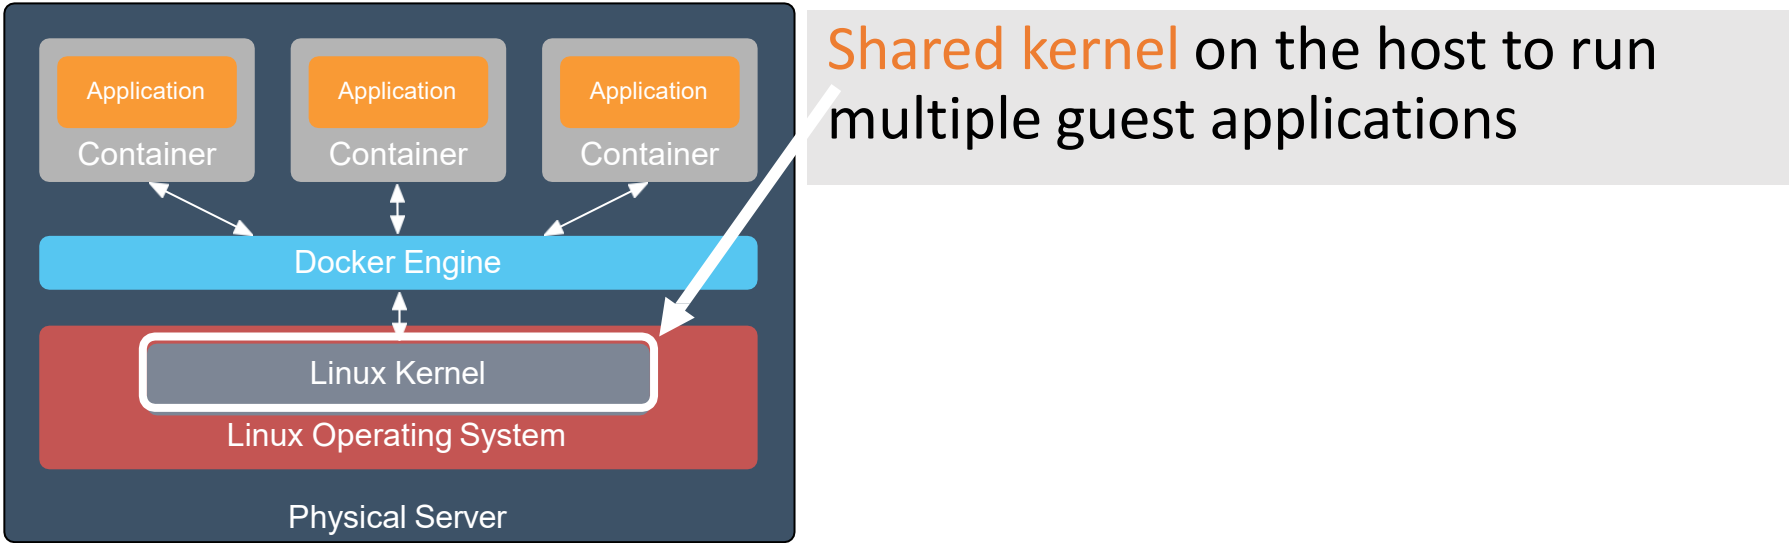

Limitations

- Each VM requires:
 - CPU allocation
 - Storage
 - RAM
 - Guest Operating System
- More VMs, more wasted resources
- Application portability not guaranteed

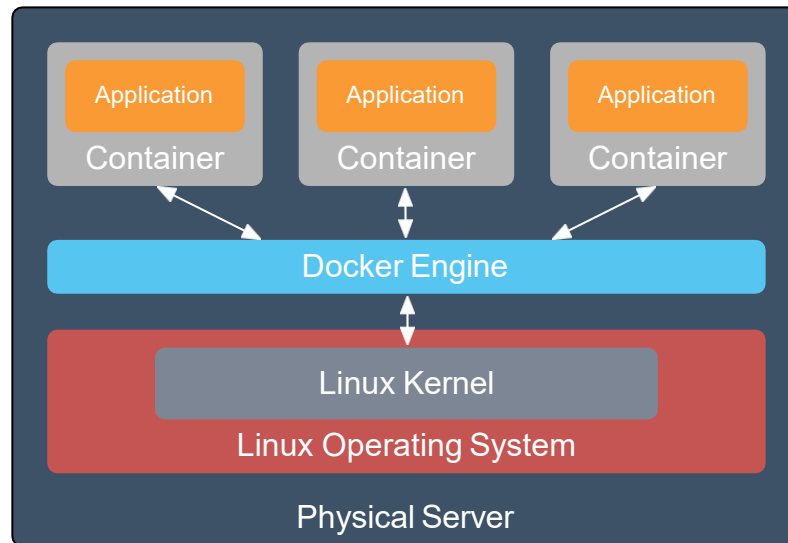

Containers



Containers



Containers - Advantages





Shared kernel on the host to run multiple guest applications

Advantages over VMs

- Containers are more lightweight
- No need to install a guest Operating System
- Less CPU, RAM, storage overhead
- More containers per machine
- Greater portability

Containers - Challenges



Shared kernel on the host to run multiple guest applications

Container Challenges

- Early Docker focused on single-node operations
- Up to user to cluster Docker hosts and manage deployment of containers on cluster
- User solves for automatic scale out of applications
- User solves for service discovery between application components (microservices)

Container – Concept of Operations



Container based virtualization

Uses the kernel on the host operating system to run multiple guest instances

- Each guest instance is a container
- Each container has its own

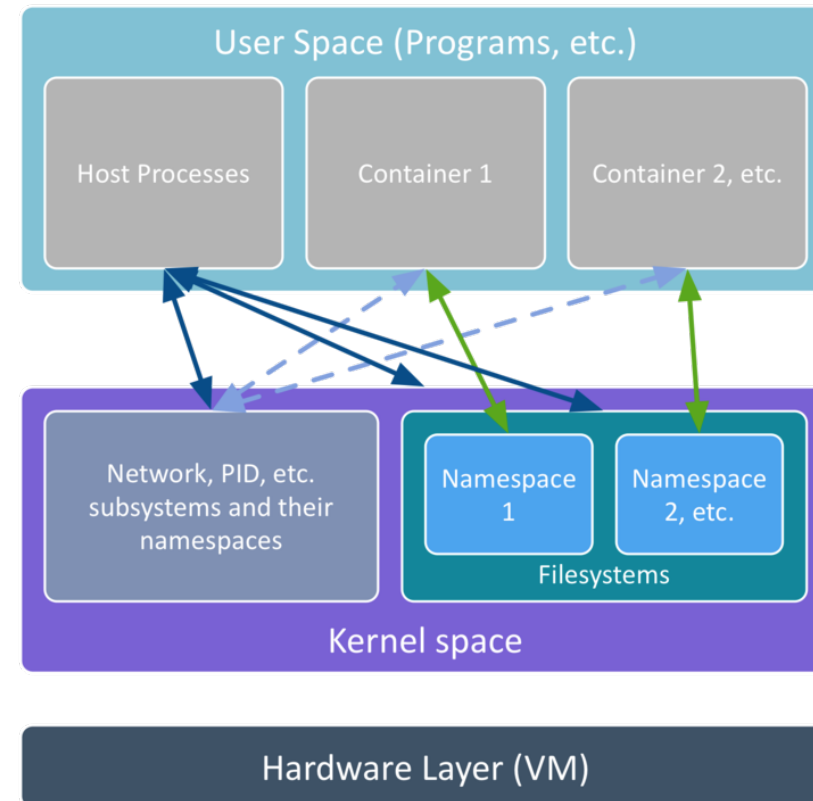
- Root filesystem
- Processes
- Memory
- Devices
- Network Ports



Isolation with Namespaces

Namespaces - Limits what a container can see (and therefore use)

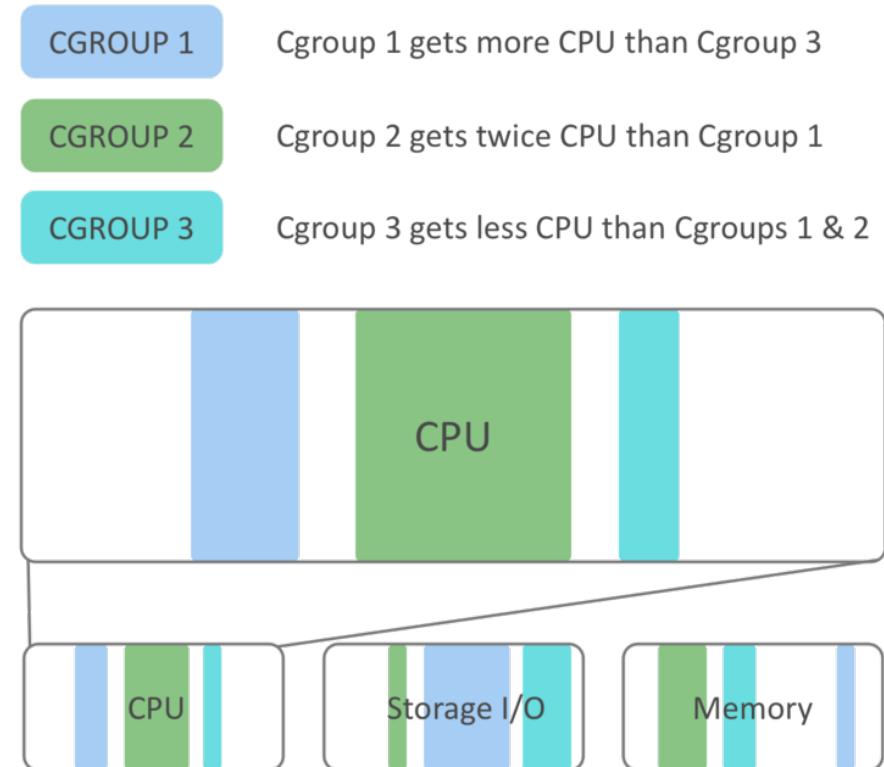
- Namespace wrap a global system resource in an abstraction layer
- Processes running in that namespace think they have their own, isolated resource
- Isolation includes:
 - Network stack
 - Process space
 - Filesystem mount points
 - etc.



Isolation with Control group (Cgroups)

Cgroups - Limits what a container can use

- Resource metering and limiting
 - CPU
 - MEM
 - Block/IO
 - Network
- Device node (/dev/*) access control



Docker resource limits

```
play@lab19:~$ docker run --help |grep -E 'cpulmem|blkio'
```

Options:

<code>--blkio-weight</code> uint16	Block IO (relative weight), between 10 and 1000, or 0 to disable (default 0)
<code>--blkio-weight-device</code> list	Block IO weight (relative device weight) (default [])
<code>--cgroup-parent</code> string	Optional parent cgroup for the container
<code>--cpu-period</code> int	Limit CPU CFS (Completely Fair Scheduler) period
<code>--cpu-quota</code> int	Limit CPU CFS (Completely Fair Scheduler) quota
<code>--cpu-rt-period</code> int	Limit CPU real-time period in microseconds
<code>--cpu-rt-runtime</code> int	Limit CPU real-time runtime in microseconds
<code>-c, --cpu-shares</code> int	CPU shares (relative weight)
<code>--cpus</code> decimal	Number of CPUs
<code>--cpuset-cpus</code> string	CPUs in which to allow execution (0-3, 0,1)
<code>--cpuset-mems</code> string	MEMs in which to allow execution (0-3, 0,1)
<code>--device-read-iops</code> list	Limit read rate (IO per second) from a device (default [])
<code>--device-write-iops</code> list	Limit write rate (IO per second) to a device (default [])

<code>-m, --memory</code> bytes	Memory limit
<code>--memory-reservation</code> bytes	Memory soft limit
<code>--memory-swap</code> bytes	Swap limit equal to memory plus swap: '-1' to enable unlimited swap
<code>--memory-swappiness</code> int	Tune container memory swappiness (0 to 100) (default -1)

Container Use Cases



DevOps



Developers

Focus on applications inside the container



Operations

Focus on orchestrating and maintaining
containers in production

Container use-cases

Development

- Allows the ability to define the entire project configuration and tear-down/recreate it easily
- Supports multiple versions of application simultaneously

Container use-cases

Development

- Allows the ability to define the entire project configuration and tear-down/recreate it easily
- Supports multiple versions of application simultaneously

Test Environments

- Same container that developers run is the container that runs in test lab and production – includes all dependencies
- Well formed API allows for automated building and testing of new containers

Container use-cases

Development

- Allows the ability to define the entire project configuration and tear-down/recreate it easily
- Supports multiple versions of application simultaneously

Test Environments

- Same container that developers run is the container that runs in test lab and production – includes all dependencies
- Well formed API allows for automated building and testing of new containers

Microservices

- Design applications as suites of services, each written in the best language for the task
- Better resource allocation
- One container per microservice vs. one VM per microservice
- Can define all interdependencies of services with templates

Questions




Container Overview



What is Docker?




Docker allows you to package an application with all of its dependencies into a standardized unit for software development.



Terminology

Image



Read only template used to create containers

Built by you or other Docker users

Stored in Docker Hub, Docker Trusted Registry or your own Registry

Terminology

Image

Read only template used to create containers

Built by you or other Docker users

Stored in Docker Hub, Docker Trusted Registry or your own Registry

Container

- Isolated application platform
- Contains everything needed to run your application
- Based on one or more images

Container Images



Lab: Access VMs



Lab: Container resource limits



Docker Image Terminology

Image

Hierarchy of files, with meta-data for how to create/run a container

Union Filesystem


- Read only template used to create containers
- Can be exported or modified to new images

Dockerfile

- Created manually or through automated processes

Container

- Stored in a Registry (Docker Hub, Docker Trusted Registry, etc.)



Docker Image Terminology

Image

Union
Filesystem

UnionFS – Used by Docker to layer images

- Not a distributed File System

Dockerfile

Container

Docker Image Terminology

Image

Union
Filesystem


Dockerfile

Container

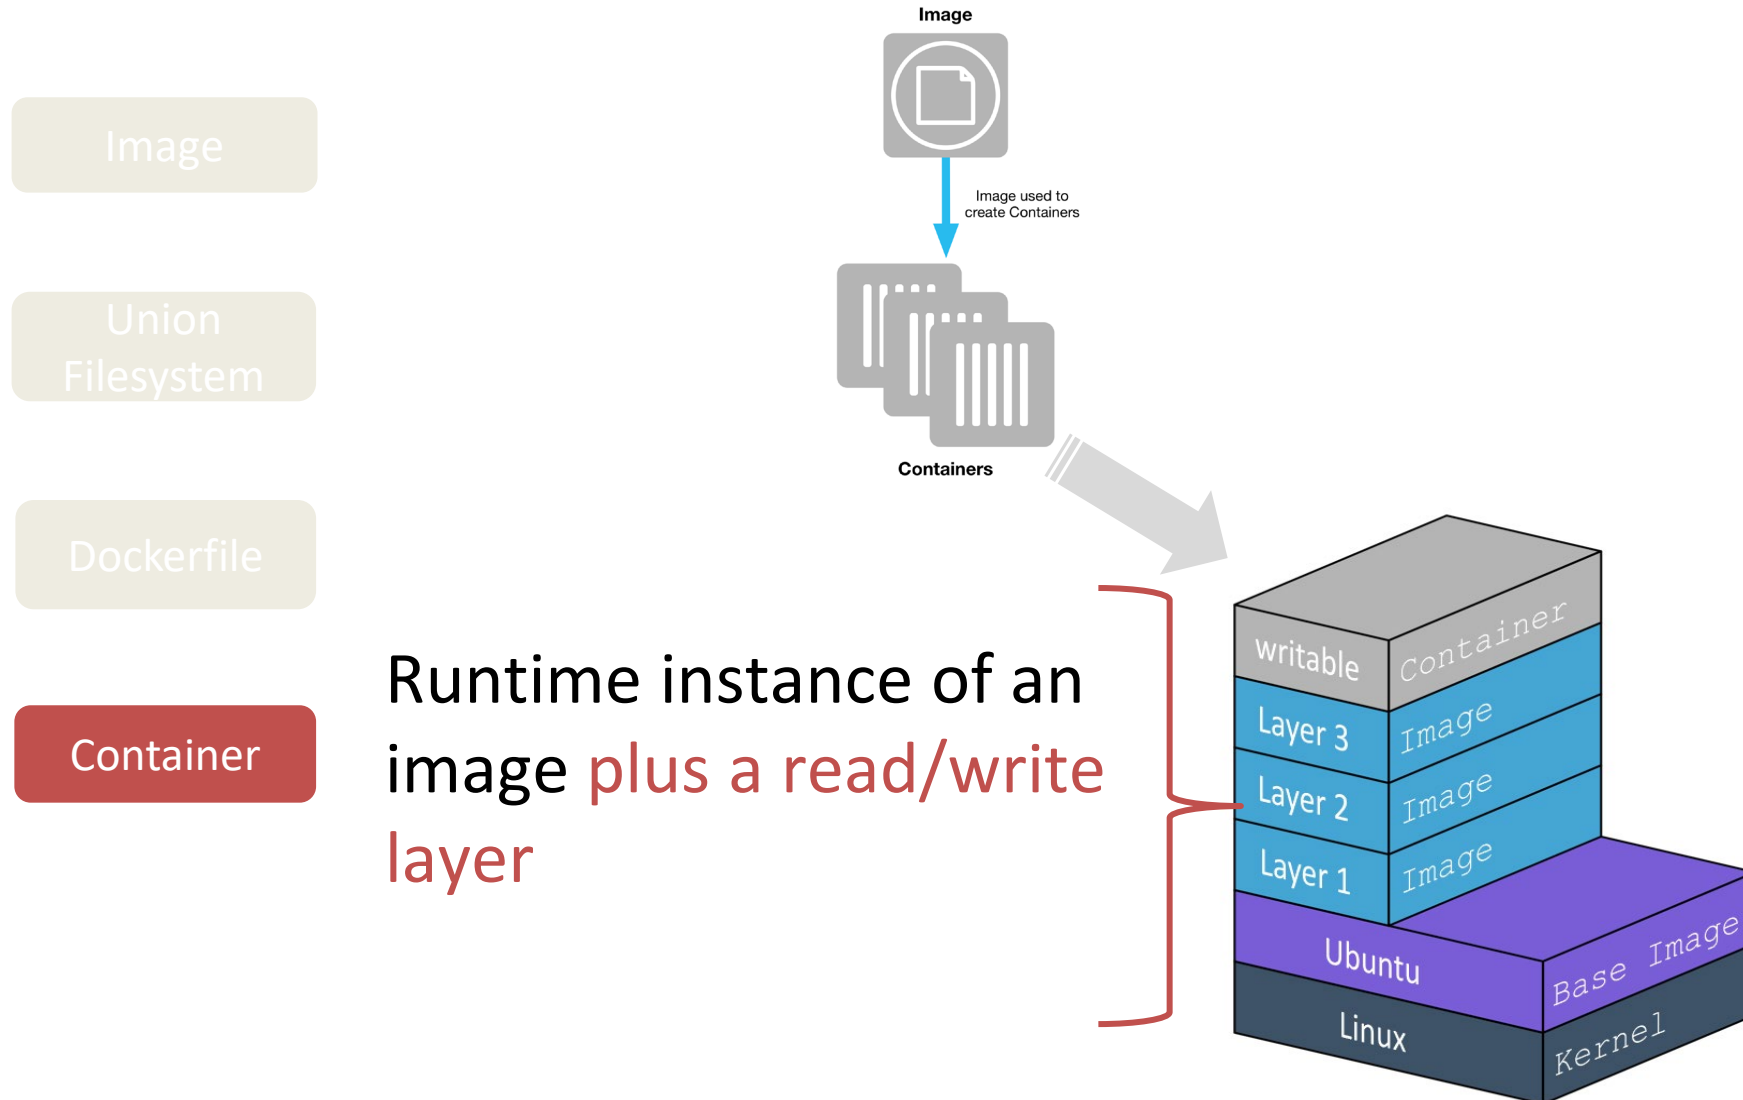
Configuration File (script) for creating images.

Defines:

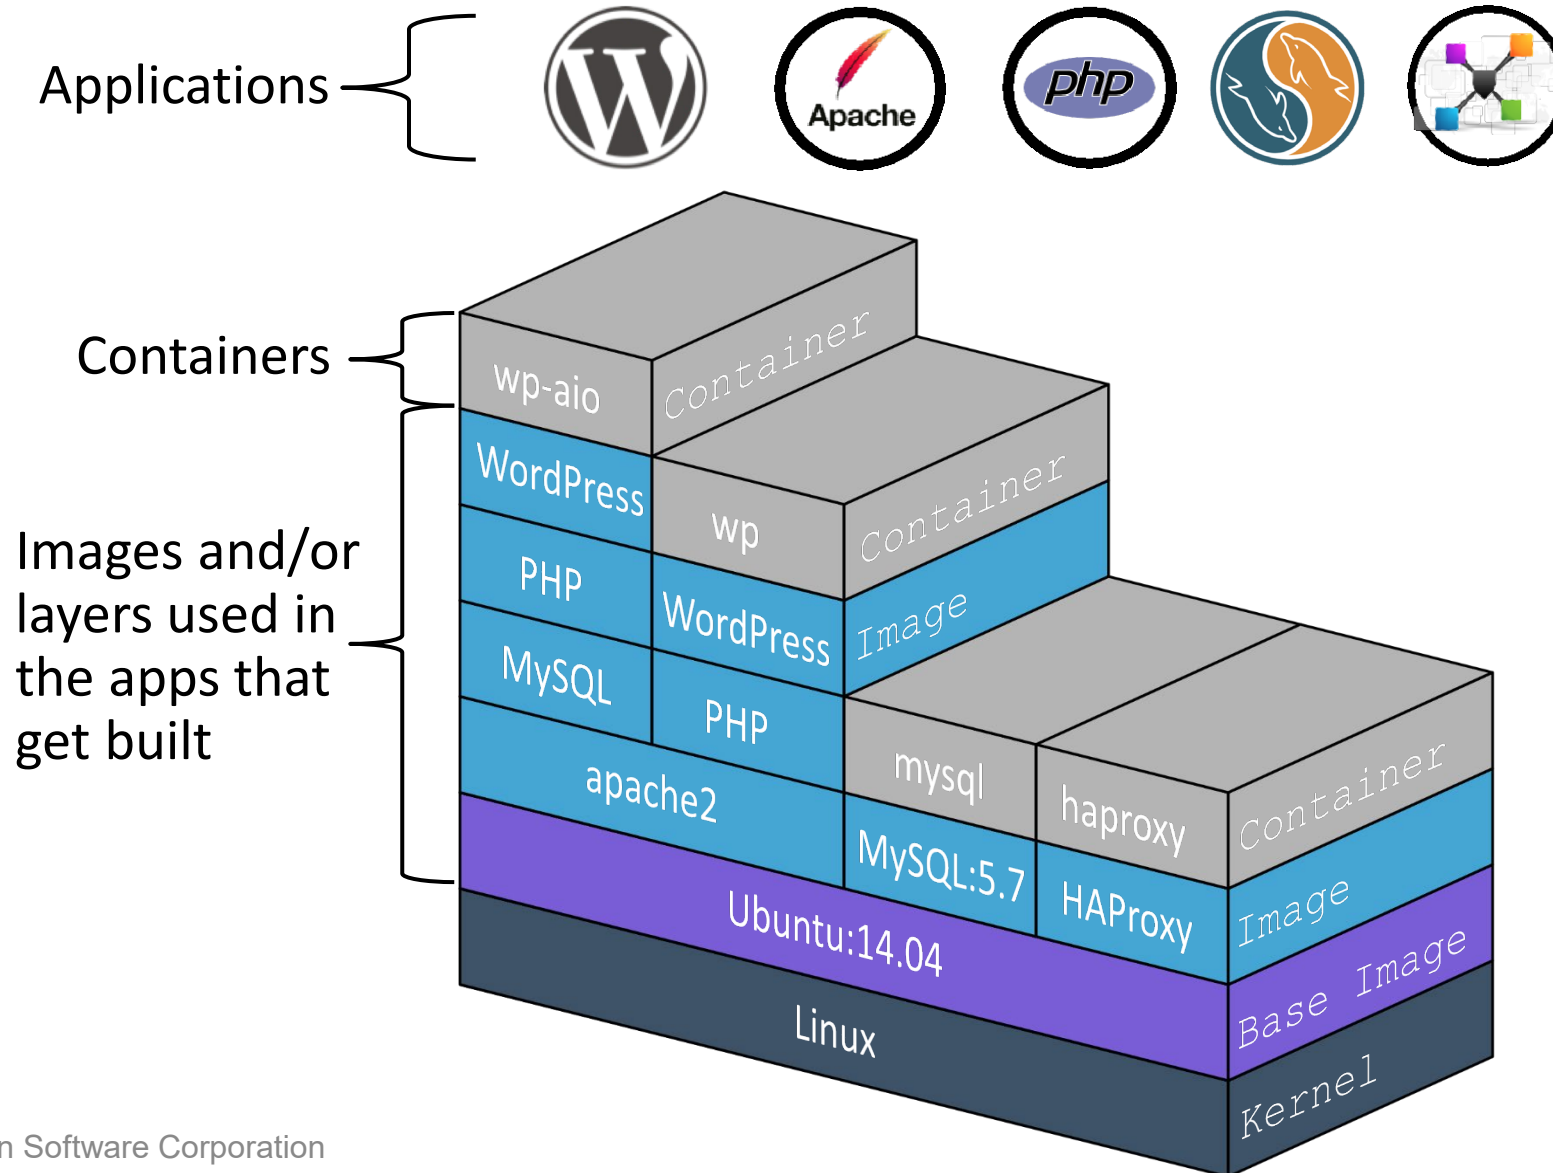
- Existing image to be the starting point
- Set of instructions to augment that image (each of which results in a new layer of the file system)
- Meta-data such as ports exposed
- The command to execute when the image is run



Docker Image Terminology



Applications, Containers, and Images



Union Files System



Image/Container Interactions

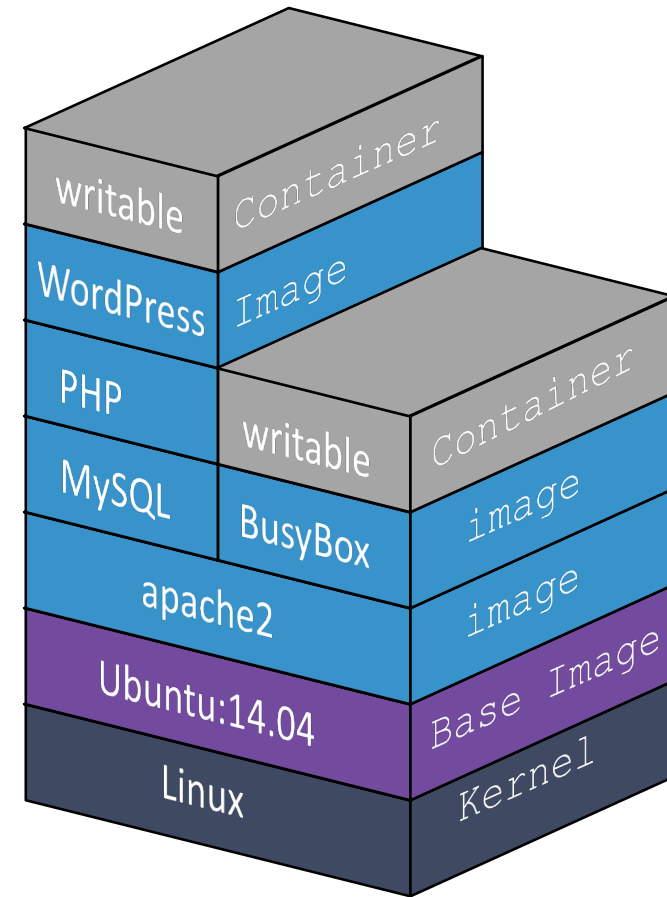
Pull an image from a Docker Registry

Run a container from an image

Add a file to a running container –
Example, the image requires an additional file called index.html

Change to a running container –
Example, the image requires an update of the index.php file

Delete files from a running container –
All containers share the same host kernel



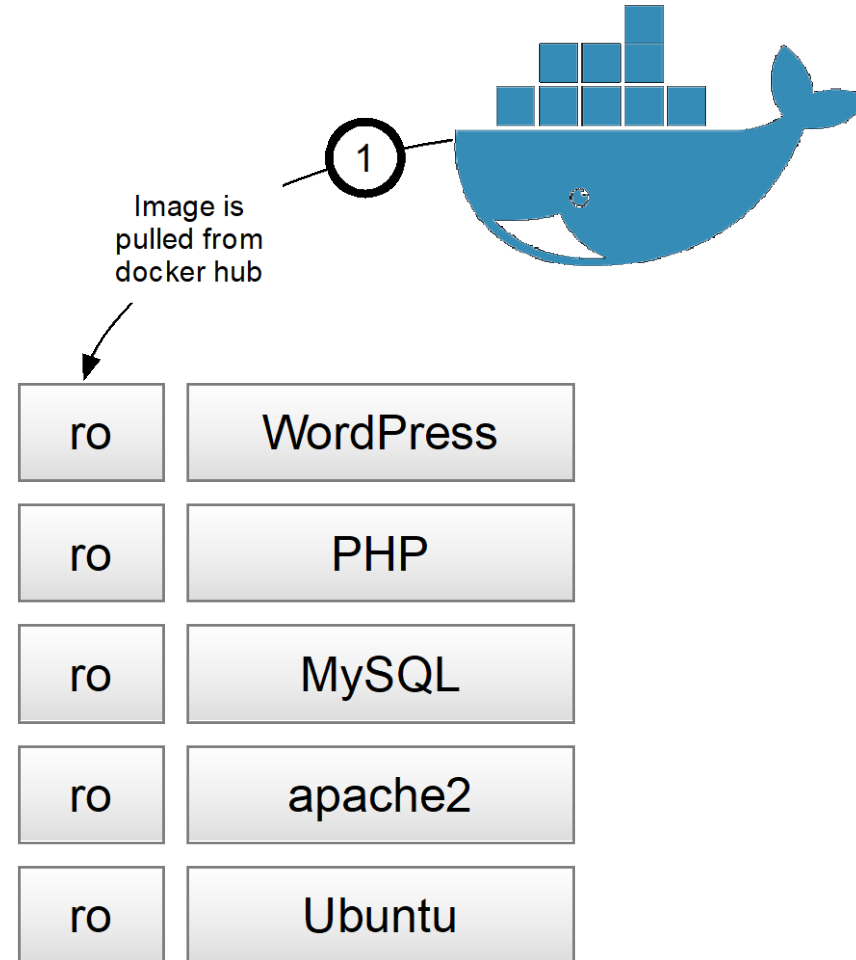
Pull the WordPress AIO

The WordPress All-In-One image is pulled from Docker Hub

```
$docker inspect wordpress-aio
```

Docker images are:

- Pulled or created
- Stored in a local image repo
- READ-ONLY
- The basis of all containers



Run the WordPress AIO

A container is made from the previously pulled WordPress AIO image

1. `$docker run wordpress-aio`
2. Container created
3. READ/WRITE layer created
4. All **ADD, CHANGE, DELETE** are committed to the READ/WRITE Layer

ro

Ubuntu

Key Takeaways

- **Images** are:
 - Highly portable and readily available
 - Reusable for many deployments
- **ADD** – ADD DATA to READ/WRITE Layer
- **CHANGE** – ADD DATA to READ/WRITE Layer
- **DELETE** – ADD DATA to READ/WRITE Layer
- **ALL** ADDs, CHANGEs, and DELETEs are copied to the READ/WRITE Layer



READ/WRITE Layer

Dockerfile best practices



Advanced Docker builds

- A great docker image is small e.g. < 20MiB
- The ideal image contains only the application*
- A great image has very few layers*
- The ideal image has 1 layer*
- The ideal build is 100% reproducible
- The ideal build is as fast as possible

*except when it doesn't



Q: When does performance matter?

Q: When does performance matter?

ALWAYS

Q: When does performance not matter?

Q: When does performance not matter?

NEVER

Typical order of operations

1. Correctness and Integrity
2. Readability and Usability
3. Performance

That said, let's take performance where it's free

Why care?

Smaller Images Mean:

- Faster builds
- Faster pulls
- Faster pushes
- Faster development
- Faster testing
- Faster deployment
- Faster feedback

Minimal base images

- scratch
- alpine (<5MiB)



Application Base Images

- redis:alpine
- mysql



Stack Base Images

- gcr.io/distroless (python, node, java, dotnet, cc)
- Python:alpine
- php:alpine
- golang



Layer review

- Images are stacks of layers
- Each Dockerfile line is a layer

```
#Layer n
FROM nixos/nix

#Layer n+1
ENV build=dev

#Layer n+2
CMD /app/bin/start

#Layer n+3
COPY app /app

#Layer n+4
RUN /app/bin/build.sh
```


Add small layers

Might be small, might be HUGE

```
FROM alpine  
COPY . .
```

Probably smaller

```
FROM alpine  
COPY app /app
```

Independent RUNs at top

RUNs are cacheable

```
FROM alpine  
RUN dd if=/dev/random of=/tmp/seed bs=1 count=1  
# More stuff ...
```

Dependent RUNs immediately after the file they depend on

```
#Layer n-  
FROM nixos/nix-  
-  
#Layer n+1-  
ENV build=dev-  
-  
#Layer n+2-  
CMD /app/bin/start-  
-  
#Layer n+3-  
COPY app /app-  
-  
#Layer n+4-  
RUN /app/bin/build.sh-  
-
```

Chain your RUNs

Chained commands happen in the same layer

```
FROM node:6-alpine
CMD ["node", "/www/src/index.js"]
COPY package.json /www/
RUN apk add --update git && \
    npm --prefix=/www install && \
    apk del git && \
    rm -rf /var/cache/apk/*
COPY src/ /www/src
```

Cleanup your RUNs

This build requires git, but the runtime doesn't. Let's cleanup after the build.

```
FROM node:6-alpine
CMD ["node", "/www/src/index.js"]
COPY package.json /www/
RUN apk add --update git && \
    npm --prefix=/www install && \
    apk del git && \
    rm -rf /var/cache/apk/*
COPY src/ /www/src
```

Add layers sparingly

Each line should have a good reason to exist

Common extraneous layers:

RUN followed by RUN

MAINTAINER <in+version+control@example.com>

EXPOSE 80

WORKDIR

Sometimes extraneous:

VOLUME /app/data

Order matters

Anti-Example:

```
COPY . /tmp/↵  
RUN pip install --requirement /tmp/requirements.txt↵  
↵
```

Not bad?:

```
COPY requirements.txt /tmp/↵  
RUN pip install --requirement /tmp/requirements.txt↵  
COPY . /tmp/↵  
↵
```

Order matters

Good:

```
COPY LICENSE /tmp/ ·↵  
COPY requirements.txt /tmp/↵  
RUN pip install --requirement /tmp/requirements.txt↵  
COPY *.py /tmp/↵  
↵
```

Best:

```
COPY LICENSE /tmp/ ·↵  
COPY requirements.txt /tmp/↵  
RUN pip install --requirement /tmp/requirements.txt↵  
COPY constants.py /tmp/↵  
COPY main.py /tmp/↵  
↵
```


.dockerignore - reduce build context

you can use .dockerignore in root of build context to exclude files from the build

.dockerignore makes these better:

```
COPY ..  
COPY src/ /app/src
```

When simple, explicit is still preferable:

```
COPY foo /app/foo  
COPY bar /app/bar
```

When to **not** minimize an image

When it impacts usability or readability:

- run a debugger inside the container (copy in the debugger dependencies)
- run a profiler inside the container (copy in the profiler dependencies)
- run a supporting application inside the container (e.g. redis-cli)
- open a shell inside a container

Multi-stage Docker builds

- Multiple FROM statements can be used to decompose a docker build into “stages”
- Conditionally build a single stage, or all stages

Common multi-stage patterns

- Create a “build” layer that performs the actual compilation of your application
- Create a “test” layer that contains canned dummy data for your application to consume or expose for testing
- Create a “production” layer that is stripped down to only a lightweight base image (like alpine) and your binary

Benefits of multi-stage builds

Portability

- Utilizing a "build" layer in your Dockerfile allows a docker build to be ported to any build system (Jenkins, TravisCI, CircleCI, etc) that can build a container. No consideration has to be made as to whether the build system "supports" builds with your preferred runtime.

Security

- By only including explicitly what you want in the final image, you decrease your container attack surface

Simplicity

- Removes any requirement to have "test" and "production" Dockerfiles to facilitate testing.
- The number and size of layers in intermediate stages does not factor into the final deployable image.

Multi-stage build example

Create a named “build” stage

Build stage contains secrets that should not exist in the final image

Create an unnamed runtime stage

Copy artifacts from the “build” stage to the runtime stage

```
FROM node:8-alpine AS build
COPY .npmrc /home/user/.npmrc
COPY package.json /app/
RUN npm --prefix=/app install
COPY src /app/src

FROM node:8-alpine
CMD ["node", "/app/src/index.js"]
COPY --from=build /app /app
```

Questions



Lab: Multi-stage Dockerfiles



Lab: Analyzing images

